

Leveraging Word Embedding from Macro and Micro View to Boost Performance for Semantic Textual Similarity

June 10, 2016

Outline

- Task Definition
- Our Systems
 - Preprocess
 - Traditional NLP Feature Engineering
 - Word Embedding Feature Engineering
- Experiments
- Results
- Conclusion

Task Definition

Definition (Semantic Textual Similarity)

Input: given two sentences

Output: similarity score([0,5])

Gold Standard: human judgements

Evaluation: Pearson correlation

Example

{ The bird is bathing in the sink.
Birdie is washing itself in the water basin. (sys: ? / gs: 5.0)

{ The woman is playing the violin.
The young lady enjoys listening to the guitar. (sys: ? / gs: 1.0)

Our Systems

Traditional NLP Feature Engineering


- 1 String-Based Similarity
- 2 Machine Translation Similarity
- 3 Corpus-based Features
- 4 Alignment Measures

Word Embedding Feature Engineering

- 1 Word Centroid Distance
- 2 Word Mover's Distance

Preprocess

- 1 To Formal Writing (Search and Replace)
e.g., doesn't → does not
- 2 Lemmatize (NLTK and Stanford CoreNLP)
e.g., was → be
- 3 Parse (Stanford CoreNLP)



Traditional NLP Feature Engineering

- 1 String-Based Similarity
- 2 Machine Translation Similarity
- 3 Corpus-based Features
- 4 Alignment Measures

String-Based Similarity

- Length Features (len):**

$$|A|, |B|, |A - B|, |B - A|, |A \cup B|, |A \cap B|, \frac{|A - B|}{|B|}, \frac{|B - A|}{|A|}$$

- Syntactic Features (pos):**

$$|A_{pos}|, |B_{pos}|, |A_{pos} - B_{pos}|, |B_{pos} - A_{pos}|$$

$$|A_{pos} \cup B_{pos}|, |A_{pos} \cap B_{pos}|, \frac{|A_{pos} - B_{pos}|}{|B_{pos}|}, \frac{|B_{pos} - A_{pos}|}{|A_{pos}|}$$

- Longest Common Sequence (lcs):**

$$\frac{|lcs(A, B)|}{\min(|A|, |B|)}$$

String-Based Similarity

- **Ngrams Overlap Features (ngram):**

- 1 word level (original and lemmatized) / character level.
- 2 $n = \{1, 2, 3\}$ are used for the word level.
- 3 $n = \{2, 3, 4\}$ are used for the character level.

- **Named Entities Features (ner):**

location, organization, data, money, person, time, percent

Machine Translation Similarity

Machine Translation Similarity

1. Viewed as one input and one output of a MT system.
2. MT measures (i.e., WER, TER, PER, NIST, ROUGE-L, GTM-1)
3. Two strategies (i.e., average and concatenate)

Corpus-based Features

- **WordNet Rank Features (wordnet):**

- 1 normalized ranking (sentence) vector.
- 2 sentence vector distance: cosine, manhattan, Euclidean, Jaccard.

- **Vector Space Sentence Similarity (lsa):**

- 1 Latent Semantic Analysis(LSA)
- 2 New York Times Annotated Corpus(NYT) / Wikipedia
- 3 convert to sentence level: sum up / use idf to weigh each word vector.

Alignment Measures

$$\left\{ \begin{array}{l} \frac{12 \text{ killed in bus accident in Pakistan.}}{10 \text{ killed in road accident in NW Pakistan.}} \end{array} \right. \quad (\text{sys: } (2/3)*5 / \text{gs: } 3.2)$$

- Global Alignment Features:**

$$\text{sim}(S_1, S_2) = \frac{n_a(S_1) + n_a(S_2)}{n(S_1) + n(S_2)}$$

- POS-Specific Alignment Features:**

calculate the aligned words proportion specifically according to POS tag(i.e., noun, verb, adjective, adverb).

Word Embedding Feature Engineering

- 1 Word Centroid Distance
- 2 Word Mover's Distance




Figure: An illustration of the word centroid distance.

Word Embedding Feature Engineering

- 1 Word Centroid Distance
- 2 Word Mover's Distance

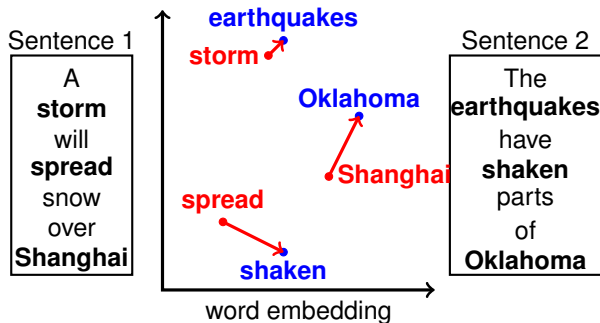


Figure: An illustration of the word mover's distance.

Word Embedding Feature Engineering

Word Embedding

- 1 word2vec
- 2 GloVe
- 3 C&W
- 4 Wiki

Distance Measurements

- 1 Cosine Distance
- 2 Manhattan Distance
- 3 Euclidean Distance
- 4 Pearson coefficient
- 5 Spearman coefficient
- 6 Kendall tau coefficient

Datasets

Training Set			Test Set		
Dataset	Input	Gold	Dataset	Input	Gold
MSRpar	1500	1500	answers-answers	1572	254
SMTeuroparl	1193	1193	plagiarism	1271	230
headlines*	3000	2250	headlines*	1498	249
SMTnews	399	399	postediting	3287	244
MSRvid	1500	1500	question-question	1555	209
OnWN	2061	2061	-	-	-
FNWN	189	189	-	-	-
images	2250	1500	-	-	-
deft-forum	450	450	-	-	-
deft-news	300	300	-	-	-
tweet-news	750	750	-	-	-
answers-forums	1500	375	-	-	-
answers-students	1500	750	-	-	-
belief	2000	375	-	-	-
All	19092	13592	All	9183	1186

Table: The statistics of all datasets for STS task.

Questions

Q₁ Supervised Model?

A: Learning Algorithm: SVM? RF? GB?

Q₂ Difference between Training Data and Test Data?

A: Training Data: All? Selected Data?

Q₃ Efficient Feature Set?

A: Feature Set: ? ?

Q₁ : Learning Algorithm

Regression	belief	answers -students	headlines	images	answers -forums	Weighted Mean
SVR(c=1)	0.7413	0.7359	0.8168	0.8660	0.7400	0.7898
RF(n=40)	0.7466	0.7100	0.8200	0.8534	0.7398	0.7816
GB(n=140)	0.7655	0.7484	0.8439	0.8791	0.7469	0.8080
DLSCU-S1	0.7491	0.7725	0.8250	0.8644	0.7390	0.8015

Table: Results of different algorithm on STS 2015 test data.

Q₂ : Training Data

Measurements

- 1 source
- 2 average length of sentences
- 3 word mover's distance

Training Data for STS 2016 test data

- 1 **headlines:** headlines
- 2 **answers-answers, question-question:**
belief, debt-forums, answers-students, answers-forums
- 3 **postediting:** SMTeuopar, MSRpar
- 4 **plagiarism:** onWN, FNWN

Q₃ : Feature Selection

Feature		belief	answers -students	headlines	image	answers -forums
String-based	len	-	-	√	√	-
	pos	-	-	√	-	√
	lcs	-	-	√	√	√
	ngram	-	√	√	√	√
	ner	-	√	√	-	-
Machine Translation	average	√	-	√	√	-
	concat	-	-	-	-	-
Corpus-based	wordnet	√	√	√	-	√
	lsa	-	√	√	√	√
Alignment	global	-	√	√	√	√
	specific	√	√	√	√	√
Word Centroid Distance	word2vec	√	√	√	√	√
	glove	-	-	-	-	-
	turian's	√	√	√	√	-
Word Mover's Distance	wmd	√	√	√	√	√
Our Results		0.7835	0.7713	0.8455	0.8808	0.7636
Best Scores		0.7717	0.7879	0.8417	0.8713	0.7390

Table: Results of feature selection experiments on STS 2015 test data.

Setups

- U-SEVEN:
 - 1 longest common sequence
 - 2 alignment feature
 - 3 corpus-based feature
 - 4 word centroid distance from four word embedding.
cosine distance, Pearson coefficient, Spearman coefficient.
- S1-All
 - 1 all the training datasets
 - 2 regression model: GB($n=140$)
 - 3 feature selection: hill climbing
- S2
 - 1 selected training datasets
 - 2 regression model: GB($n=140$)
 - 3 feature selection: hill climbing

Results

Dataset	Runs			Best Score
	U-SEVEN	S1-All	S2	
answers-answers	0.4774	0.5697	0.5715	0.6923
plagiarism	0.8301	0.8250	0.7733	0.8413
headlines	0.7668	0.8121	0.7903	0.8274
postediting	0.8423	0.8234	0.7496	0.8669
question-question	0.7191	0.7311	0.6763	0.7470
weighted mean	0.7242	0.7507	0.7116	0.7780

Table: The results of our three runs on STS 2016 test datasets.

Results

{	You should do it.	(sys: 4.0 / gs: 1.0)
{	You can do it, too.	
{	It's pretty much up to you.	(sys: 3.2 / gs: 0.0)
{	It's much better to ask.	

Conclusion

- 1 The difference between top system and our best system is about 2.8%
- 2 Future work: Deep Learning